



Real-Time Systems

Challenges using Linux to implement Real-Time Services

Dr. Sam Siewert

Electrical, Computer and Energy Engineering

Embedded Systems Engineering Program

Overview

- Review of Linux approach compared to RTOS and Cyclic Executive
- Linux Kernel Space – Option #2 (Why not just use an RTOS!)
- Linux User Space – Option #1 (Has advantages, but high overhead)
- Why User Space is good
- Why User Space is bad
- How to use a combination of User and Kernel space with RT configuration

3 Approaches to RT Systems with Software Services

Cyclic Exec

E.g. Shuttle Flight Software, Network elements, process control, Ada83/95 CE

Custom, Deeply Embedded Systems

Low over-head, purpose-built

Costly to develop

RTOS

VxWorks, QNX, ThreadX, TI RTOS, FreeRTOS, Zephyr, Nucleus, RTEMS, etc.

Embedded Systems,
Scalable and Portable

Medium overhead, quick to market

License costs

OS + RT POSIX

RT Linux, Solaris, LynxOS, FSM Labs, Concurrent, RTAI, Linux Foundation RT, etc.

RT Services for Scalable Apps and Systems

Highest overhead, mixed RT + SRT + BE

Maintenance costs

Linux User Space vs. Kernel Space

- Linux has 2 address spaces – user and kernel
 - Kernel space directly interfaces to hardware and manages all resources
 - Provides protected address spaces for user processes (that can have threads mapped onto kernel tasks)
 - Provides security and stability
 - Driver and kernel modules have different development, debug, and test strategies!
 - User space protects the system from errors and provides security & stability
 - All applications interfacing to hardware must go through a driver and/or system call
 - All applications run in within a protected address and resource domain
 - Can we have both? (cake and eat it too?) – with care
 - Sub-millisecond time critical, direct hardware interfaces must be in a driver kernel module
 - Millisecond or less predictable response (soft RT and BE) services can operate in user space
 - Carefully configure a Linux platform (real-time distro) and allocate services to both user space and kernel space based on requirements and timing constraints (carefully!) – **TEST it!**

AMP Fixed Priority RM Policy on Linux

- Linux is SMP by default on Multi-core systems
 - We emulate AMP with Pthread Affinity
 - We implement Fixed Priority RM Policy with SCHED_FIFO
 - Check IRQ balance ([irqbalance](#)) and consider masking off RT cores
- We need periodic source of interrupts, driver, and sequencing (**sub-millisecond services** must be implemented in **kernel space** in a driver module)
- Without a periodic source of interrupt, we must sequence operations with timers
- User space RT services are possible, but accuracy and precision impacted by **User to Kernel system call overhead** (accurate to 1 to 10 milliseconds at best!)

Linux Kernel Patches and Configuration for Real-Time

- Linux is not configured for Real-Time in most distributions!
 - Notable exceptions: Automotive Grade Linux, Concurrent RedHawk Linux, Wind River Linux, etc.
 - We can purchase support for pre-configured RT Linux distro or roll our own
- Key patches recommended (not used in this course for simplicity)
 - RT Preempt patch for Kernel
 - RT Mutex priority inheritance semaphore patch (or use of FUTEX semaphores)
 - High Resolution Timers (enable in user space – TSC on x86/x64, STC on ARM) or keep time in kernel space
- Key configuration updates (not used for simplicity)
 - run levels – make your own for your services!
 - remove unused services, pair down features to just those required!

Avoid Large “Unknown” Libraries in Linux

- Attraction to an OS like Linux is code and platform reuse!
 - Up and running quickly, working on drivers and applications
 - Safe for multiple developers
 - Stable with security features
- Challenge is that short of reviewing all source code use, we don't know what is really going on!
- Pair down the system to essentials and use custom [RT run-level](#)
 - E.g. run-level “4” is for “special purposes”
- If large libraries are used – e.g. ROS, OpenCV, etc., then review source and test!!

What we do in ECEE 5315 – 5318 Courses Series

- We use the Raspberry Pi for convenience and simplicity
- We learn HRT, SRT, and Best Effort service theory! – useful for CE, RTOS or Linux
 - We practice implementation of RT Services as POSIX SCHED_FIFO threads
 - We work around accuracy and precision limitations in user space (millisecond accuracy is sufficient for our learning objectives – 30 Hz or less)
 - We learn about Linux kernel space options, RTOS alternatives, ARM M Series or R Series Cyclic Executive options
- At the end, as a new RT theorist and practitioner, you will decide – CE, RTOS or Linux + RT configuration and extensions?
- Optional Textbook – RTECS with Linux and RTOS covers both options
- When in doubt – go with RTOS, CE or Hardware HRT Services!

How many systems are HRT?
(loss of life, property, or \$\$\$)

- Civil aviation flight control – YES!
- Medical ICU equipment – YES!
- Public transportation – YES!
- Energy generation, distribution and monitoring – YES!
- Self-driving car, Autonomous Air taxi – YES!

But, many things are not ...
(loss of quality, not system, low or no collateral damage – safety judgement)

Robotics? - Maybe (size, power)
Digital Media – SRT QoS
Games – SRT QoS
UAV/UAS – Some HRT (autopilot), some SRT (navigation and imaging)
Driver's assistant – navigation, warnings and specific functions

Copyright © 2019 University of Colorado

